



TRƯỜNG ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

Distributed hash table (DHT)

Lecturer: Thanh-Chung Dao

Slides by Viet-Trung Tran

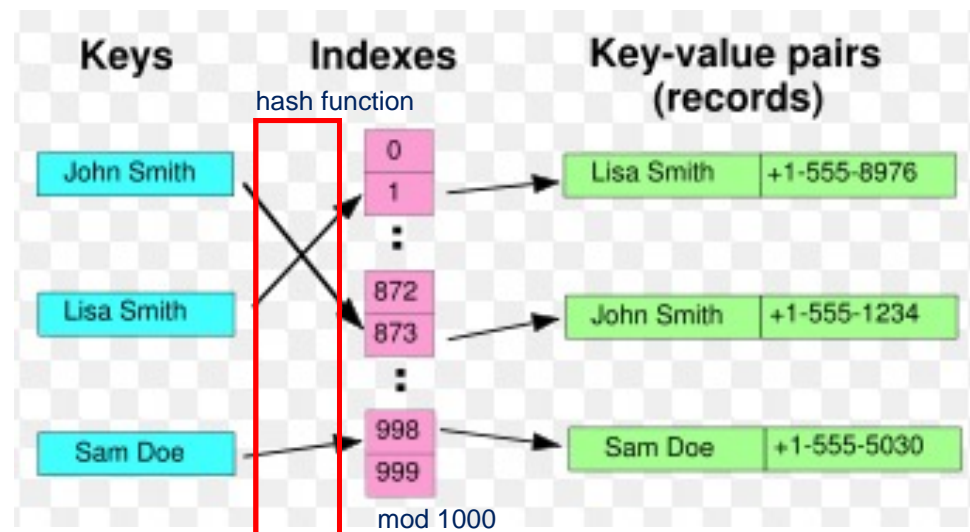
School of Information and Communication Technology

Outline

- Hashing
- Distributed Hash Table
- Chord

A Hash Table (hash map)

- A data structure implements an associative array that can map keys to values.
 - searching and insertions are $O(1)$ in the worse case
- Uses a hash function to compute an index into an array of buckets or slots from which the correct value can be found.
 - $\text{index} = f(\text{key}, \text{array_size})$



Hash functions

- Crucial for good hash table performance
- Can be difficult to achieve
 - WANTED: uniform distribution of hash values
 - A non-uniform distribution increases the number of collisions and the cost of resolving them


Hashing for partitioning usecase

- Objective
 - Given document X , choose one of k servers to use
- Eg. using modulo hashing
 - Number servers $1..k$
 - Place X on server $i = (X \bmod k)$
 - Problem? Data may not be uniformly distributed
 - Place X on server $i = \text{hash}(X) \bmod k$
- Problem?
 - What happens if a server fails or joins ($k \rightarrow k \pm 1$)?
 - What is different clients has different estimate of k ?
 - Answer: All entries get remapped to new nodes!


Distributed hash table (DHT)

- Distributed Hash Table (DHT) is similar to hash table but spread across many hosts
- Interface
 - insert(key, value)
 - lookup(key)
- Every DHT node supports a single operation:
 - Given key as input; route messages to node holding key

DHT: basic idea




DHT: basic idea




Neighboring nodes are “connected” at the application-level⁸

DHT: basic idea




Operation: take *key* as input; route messages to node holding *key*

DHT: basic idea



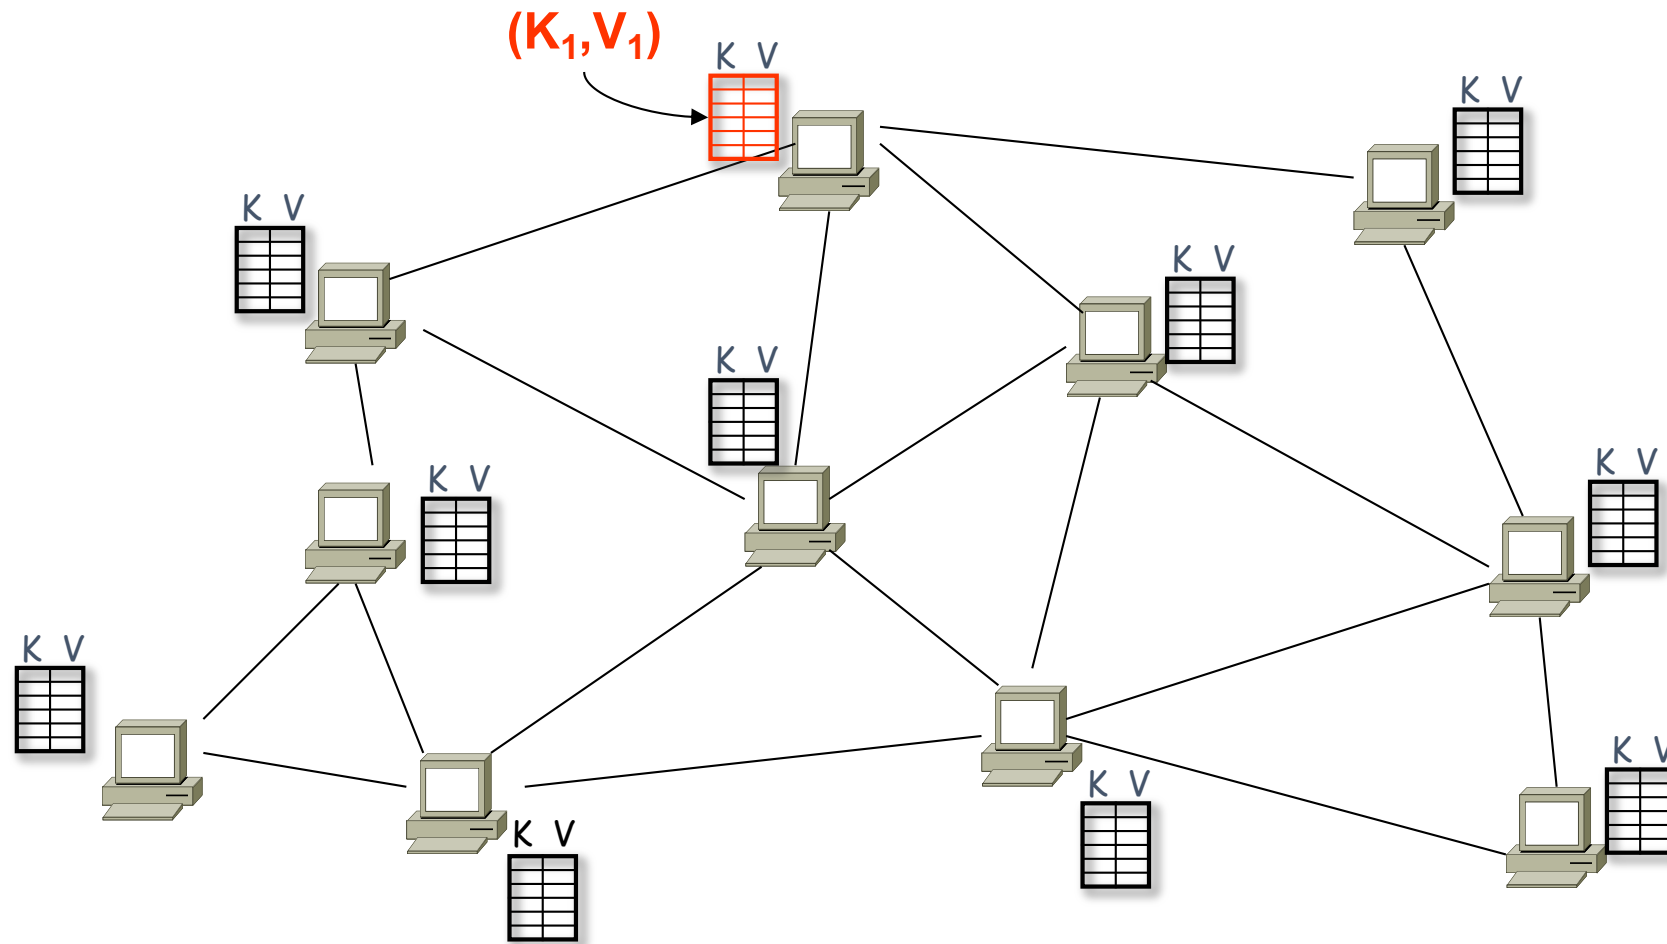
Operation: take *key* as input; route messages to node holding *key*

DHT: basic idea




Operation: take *key* as input; route messages to node holding *key*

DHT: basic idea



Operation: take *key* as input; route messages to node holding *key*

DHT: basic idea



Operation: take *key* as input; route messages to node holding *key*

How to design a DHT?

- State Assignment
 - What “(key, value) tables” does a node store?
- Network Topology
 - How does a node select its neighbors?
- Routing Algorithm:
 - Which neighbor to pick while routing to a destination?
- Various DHT algorithms make different choices
 - CAN, Chord, Pastry, Tapestry, Plaxton, Viceroy, Kademlia, Skipnet, Symphony, Koorde, Apocrypha, Land, ORDI ...



Chord: A scalable peer-to-peer look-up protocol for internet applications

Credit: University of California, berkely and Max planck institute

Outline

- What is Chord?
- Consistent Hashing
- A Simple Key Lookup Algorithm
- Scalable Key Lookup Algorithm
- Node Joins and Stabilization
- Node Failures

What is Chord?

- In short: a peer-to-peer lookup system
- Given a key (data item), it maps the key onto a node (peer) map key vào trong một node trong hệ thống cluster
- Uses consistent hashing to assign keys to nodes . sử dụng hàm băm để gán 1 key vào trong 1 node
- Solves the problem of locating key in a collection of distributed nodes.
- Maintains routing information with frequent node arrivals and departures

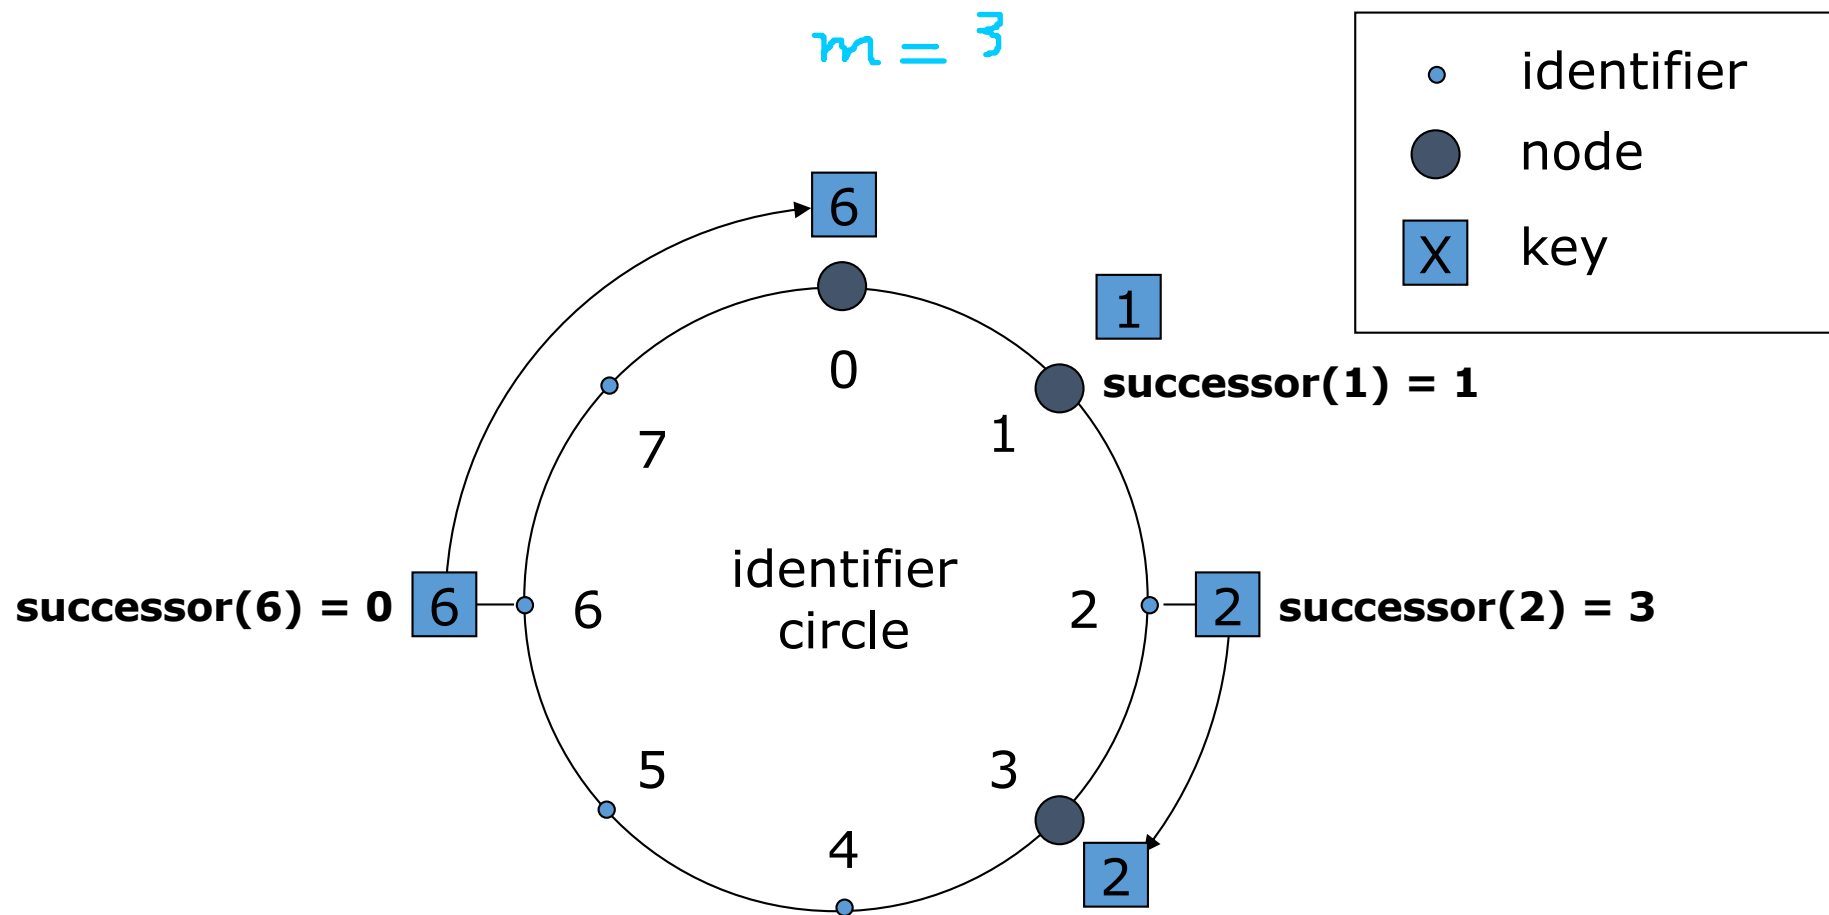
Consistent hashing

- Consistent hash function assigns each node and key an m-bit identifier. định danh m-bit 2^m identifier
- SHA-1 is used as a base hash function.
- A node's identifier is defined by **hashing the node's IP address**.
- A key identifier is produced by hashing the key (chord doesn't define this. Depends on the application).
 - $ID(\text{node}) = \text{hash}(\text{IP}, \text{Port})$
 - $ID(\text{key}) = \text{hash}(\text{key})$

Consistent hashing

- In an m -bit identifier space, there are 2^m identifiers.
- Identifiers are ordered on an identifier circle modulo 2^m .
- The identifier ring is called Chord ring.
- Key k is assigned to the first node whose identifier is equal to or follows (the identifier of) k in the identifier space.
- This node is the successor node of key k , denoted by $\text{successor}(k)$.


Consistent hashing – Successor nodes




Consistent hashing – Join and departure

- When a node n joins the network, certain keys previously assigned to n 's successor now become assigned to n .
- When node n leaves the network, all of its assigned keys are reassigned to n 's successor.

Consistent hashing – Node join



Consistent hashing – Node departure



A Simple key lookup

- If each node knows only how to contact its current successor node on the identifier circle, all node can be visited in linear order.
- Queries for a given identifier could be passed around the circle via these successor pointers until they encounter the node that contains the key.

A Simple key lookup

- Pseudo code for finding successor:

// ask node n to find the successor of id

`n.find_successor(id)`

if ($id \in (n, \text{successor}]$)

return *successor*;


else

// forward the query around the circle

return *successor.find_successor(id)*;

A Simple key lookup

- The path taken by a query from node 8 for key 54:



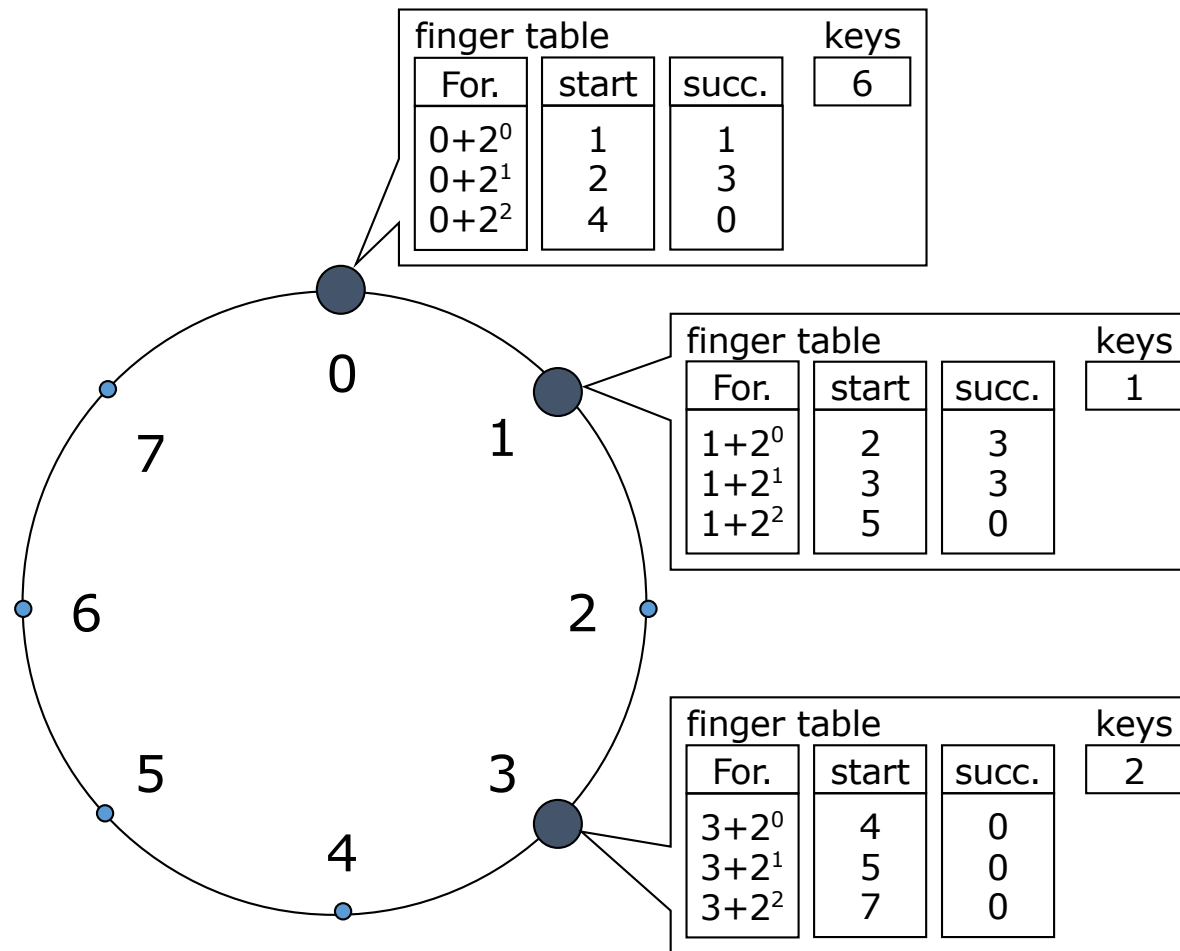
Scalable key location

- To accelerate lookups, Chord maintains additional routing information.
- This additional information is not essential for correctness, which is achieved as long as each node knows its correct successor.

Scalable key location – Finger tables

- Each node n maintains a routing table with up to m entries (which is in fact the number of bits in identifiers), called finger table.
- The i th entry in the table at node n contains the identity of the first node s that succeeds n by at least 2^{i-1} on the identifier circle.
- $s = \text{successor}(n + 2^{i-1})$. $i = 1 \rightarrow m$
- s is called the i th finger of node n , denoted by $n.\text{finger}(i)$

Scalable key location – Finger tables




Scalable key location – Finger tables

- A finger table entry includes both the Chord identifier and the IP address (and port number) of the relevant node.
- The first finger of n is the immediate successor of n on the circle.

Scalable key location – Example query

- The path a query for key 54 starting at node 8:



Scalable key location – A characteristic

- Since each node has finger entries at power of two intervals around the identifier circle, **each node can forward a query at least halfway along the remaining distance between the node and the target identifier**. From this intuition follows a theorem:
 - Theorem: With high probability, the number of nodes that must be contacted to find a successor in an N-node network is $O(\log N)$.

Node joins and stabilizations

- The most important thing is the successor pointer.
- If the successor pointer is ensured to be up to date, which is sufficient to guarantee correctness of lookups, then finger table can always be verified.
- Each node runs a “stabilization” protocol periodically in the background to update successor pointer and finger table.

Node joins and stabilizations

- “Stabilization” protocol contains 6 functions:
 - create()
 - join()
 - stabilize()
 - notify()
 - fix_fingers()
 - check_predecessor()

Node Joins – join()

- When node n first starts, it calls $n.join(n')$, where n' is any known Chord node.
- The $join()$ function asks n' to find the immediate successor of n .
- $join()$ does not make the rest of the network aware of n .

Node Joins – join()

// create a new Chord ring.

n.create()

predecessor = nil;

successor = n;

// join a Chord ring containing node n'.

n.join(n')

predecessor = nil;

successor = n'.find_successor(n);

Node joins – stabilize()

- Each time node n runs `stabilize()`, it asks its successor for its predecessor p , and decides whether p should be n 's successor instead.
- `stabilize()` notifies node n 's successor of n 's existence, giving the successor the chance to change its predecessor to n .
- The successor does this only if it knows of no closer predecessor than n .

Node joins – stabilize()

*// called periodically. verifies n's immediate
// successor, and tells the successor about n.*

n.stabilize()

x = successor.predecessor;

if (x ∈ (n, successor))

successor = x;

successor.notify(n);


// n' thinks it might be our predecessor.

n.notify(n')

if (predecessor is nil or n' ∈ (predecessor, n))

predecessor = n';

Node joins – Join and stabilization



- **n joins**
 - predecessor = nil
 - n acquires n_s as successor via some n'
- **n runs stabilize**
 - n notifies n_s being the new predecessor
 - n_s acquires n as its predecessor
- **n_p runs stabilize**
 - n_p asks n_s for its predecessor (now n)
 - n_p acquires n as its successor
 - n_p notifies n
 - n will acquire n_p as its predecessor
- **all predecessor and successor pointers are now correct**
- **fingers still need to be fixed, but old fingers will still work**

Node Joins – fix_fingers()

- Each node periodically calls fix fingers to make sure its finger table entries are correct.
- It is how new nodes initialize their finger tables
- It is how existing nodes incorporate new nodes into their finger tables.

Node Joins – fix_fingers()

// called periodically. refreshes finger table entries.

n.fix_fingers()

next = next + 1 ;

if (next > m)

next = 1 ;

finger[next] = find_successor(n + $2^{\text{next}-1}$);

// checks whether predecessor has failed.

n.check_predecessor()

if (*predecessor* has failed)

predecessor = nil;

Node failures

- Key step in failure recovery is maintaining correct successor pointers
- To help achieve this, each node maintains a successor-list of its r nearest successors on the ring. Hence, all r successors would have to simultaneously fail in order to disrupt the Chord ring.
- If node n notices that its successor has failed, it replaces it with the first live entry in the list
- Successor lists are stabilized as follows:
 - node n reconciles its list with its successor s by copying s 's successor list, removing its last entry, and prepending s to it.
 - If node n notices that its successor has failed, it replaces it with the first live entry in its successor list and reconciles its successor list with its new successor.

Chord – The math

- Each node maintains $O(\log N)$ state information and lookups needs $O(\log N)$ messages
- Every node is responsible for about K/N keys (N nodes, K keys)
- When a node joins or leaves an N -node network, only $O(K/N)$ keys change hands (and only to and from joining or leaving node)

Interesting simulation results

- Adding virtual nodes as an indirection layer can significantly improve load balance.
- The average path length is about $\frac{1}{2}(\log N)$.
- Maintaining a set of alternate nodes for each finger and route the queries by selecting the closest node according to network proximity metric improves routing latency effectively.
- Recursive lookup style is faster iterative style

Applications: Time-shared storage

- For nodes with intermittent connectivity (server only occasionally available)
- Store others' data while connected, in return having their data stored while disconnected
- Data's name can be used to identify the live Chord node (content-based routing)

Applications: Chord-based DNS

- DNS provides a lookup service
 - keys: host names values: IP addresses
- Chord could hash each host name to a key
- Chord-based DNS:
 - no special root servers
 - no manual management of routing information
 - no naming structure
 - can find objects not tied to particular machines

What is Chord? – Addressed problems

- **Load balance:** chord acts as a distributed hash function, spreading keys evenly over nodes
- **Decentralization:** chord is fully distributed, no node is more important than any other, improves robustness
- **Scalability:** logarithmic growth of lookup costs with the number of nodes in the network, even very large systems are feasible
- **Availability:** chord automatically adjusts its internal tables to ensure that the node responsible for a key can always be found
- **Flexible naming:** chord places no constraints on the structure of the keys it looks up.

Summary

- Simple, powerful protocol
- Only operation: map a key to the responsible node
- Each node maintains information about $O(\log N)$ other nodes
- Lookups via $O(\log N)$ messages
- Scales well with number of nodes
- Continues to function correctly despite even major changes of the system



Thanks for your attention!