

Computer Vision

Chapter 4: Feature detection and Image matching

1

Chapter 4 - Content

- Edge
 - Edge Detection:
 - Role of edges & general approach of edge detection
 - Image gradient
 - Edge detectors: Sobel, Prewitt, Canny detector
 - Edge Linking
- Feature extraction and matching
 - Global features
 - Local features
 - Matching and Applications

2

Edge Detection

Role of edges & general approach of edge detection

Image gradient

Edge detectors: Sobel, Prewitt, Canny detector

3

Role of edge & general approach of edge detection

What are edges/ contours ?

- Local changes in the images
- Typically occur on the boundary between different regions in an image

- Origins of edges:

Depth

Surface orientation

Surface colours

reflectance

illumination



4

Edge is important?


- (A) Cave painting at Chauvet, France, about 30,000 B.C.;
 (B) Aerial photograph of the picture of a monkey as part of the Nazca Lines geoglyphs, Peru, about 700 – 200 B.C.;
 (C) Shen Zhou (1427-1509 A.D.): Poet on a mountain top, ink on paper, China;
 (D) Line drawing by 7-year old I. Lleras (2010 A.D.).



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

5

5



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

6

6

152 Biederman




Figure 4.14
Complementary-part images. From an original intact image (left column), two complemen-

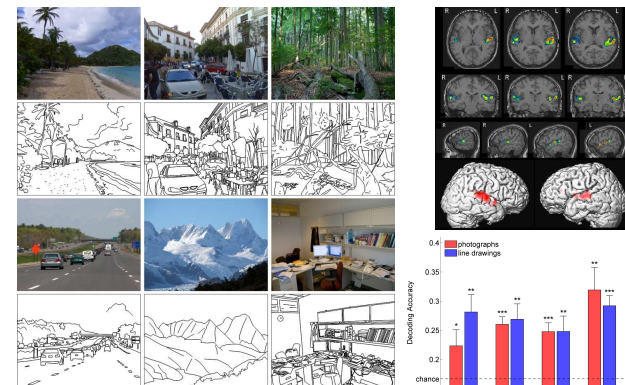
Can we recognize these objects?



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

7

7



Walther, Chai, Caddigan, Beck & Fei-Fei, Simple line drawings suffice for functional MRI decoding of natural scene categories, PNAS, 2011




SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

8

8

Edge detection

- Goal: Identify sudden changes (discontinuities) in an image
 - Intuitively, most semantic and shape information from the image can be encoded in the edges
 - More compact than pixels
- Why?
 - Extract information, recognize objects
 - Recover geometry and viewpoint



Source: J. Hayes

How to find edges?

- Intensity profile of an image is the set of intensity values taken from regularly spaced points along a line segment or multi-line path in an image





Image profile




<https://www.mathworks.com/help/images/intensity-profile-of-images.html>

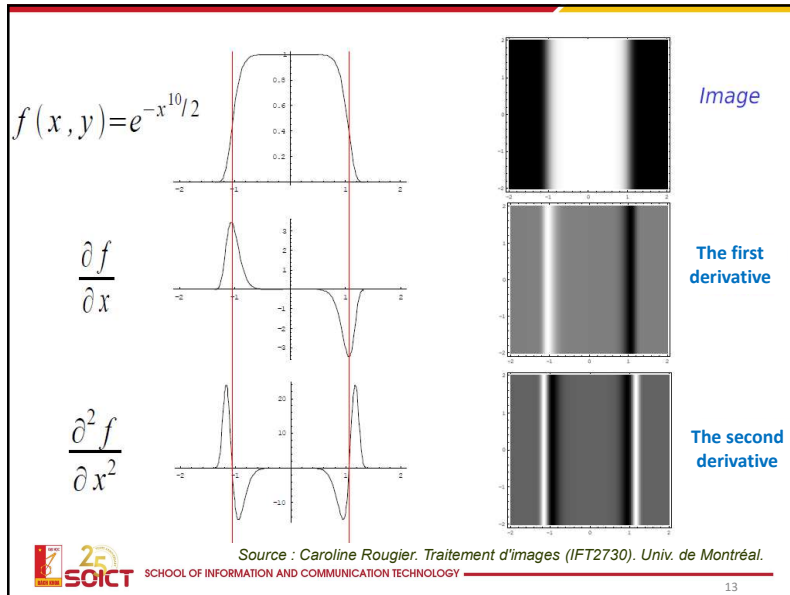
How to find edges? General approach

- An edge is a place of rapid change in the continuous image intensity function

- Extrema of 1st derivative

- Zero-crossing of 2nd derivative





13

Derivatives in 1D continuous function

- Derivative in 1D:

$$\frac{df}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x) - f(x - \Delta x)}{\Delta x} = f'(x) = f_x$$

- Discrete derivative in 1D

$$\frac{df}{dx} = \frac{f(x) - f(x-1)}{1} = f'(x)$$

$$\frac{df}{dx} = f(x) - f(x-1) = f'(x)$$

14

Types of discrete derivatives in 1D

Backward $\frac{df}{dx} = f(x) - f(x-1) = f'(x)$

Forward $\frac{df}{dx} = f(x) - f(x+1) = f'(x)$

Central $\frac{df}{dx} = f(x+1) - f(x-1) = f'(x)$

15

1D discrete derivate filters

- Backward filter:

$$\frac{df}{dx} = f(x) - f(x-1) = f'(x) \quad [0 \quad 1 \quad -1]$$

- Forward:

$$\frac{df}{dx} = f(x) - f(x+1) = f'(x) \quad [-1 \quad 1 \quad 0]$$

- Central:

$$\frac{df}{dx} = f(x+1) - f(x-1) = f'(x) \quad [1 \quad 0 \quad -1]$$

16

1D discrete derivate example

- Backward filter: $[0 \quad 1 \quad -1]$

$$f(x) = 10 \quad 15 \quad 10 \quad 10 \quad 25 \quad 20 \quad 20 \quad 20$$

$$f'(x) = 0 \quad 5 \quad -5 \quad 0 \quad 15 \quad -5 \quad 0 \quad 0$$

$$\begin{array}{l} f(x): \quad 0 \quad 0 \quad 0 \quad 0 \quad 0 \quad 50 \quad 50 \quad 50 \quad 50 \quad 50 \\ f'(x): \quad 0 \quad 0 \quad 0 \quad 0 \quad 0 \quad 50 \quad 0 \quad 0 \quad 0 \quad 0 \end{array}$$



Discrete derivate in 2D

Discrete derivate filters wrt x and y

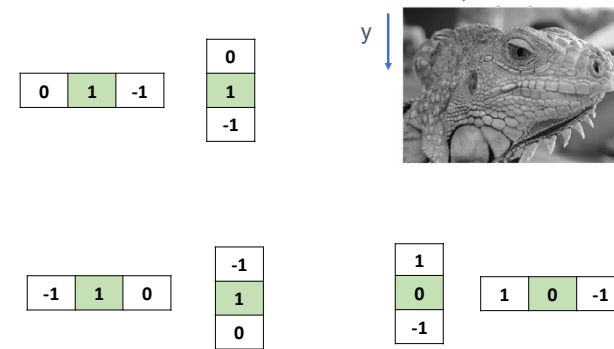


Image I (9 x 8)

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	100	100	100	100	100	100
0	0	0	100	100	100	100	100	100
0	0	0	100	100	100	100	100	100
0	0	0	100	100	100	100	100	100
0	0	0	100	100	100	100	100	100

M_x : Derivative filter wrt x

1	0	-1
---	---	----

M_y : Derivative filter wrt y

1
0
-1

$I_x = I * M_x = \text{the 1st derivative wrt x}$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	100	100	0	0	0	0
0	100	100	0	0	0	0
0	100	100	0	0	0	0
0	100	100	0	0	0	0
0	100	100	0	0	0	0
0	100	100	0	0	0	0



Image I (9 x 8)

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	100	100	100	100	100	100
0	0	0	100	100	100	100	100	100
0	0	0	100	100	100	100	100	100
0	0	0	100	100	100	100	100	100
0	0	0	100	100	100	100	100	100

M_x : Derivative filter wrt x

1	0	-1
---	---	----

M_y : Derivative filter wrt y

1
0
-1

$I_y = I * M_y = \text{the 1st derivative wrt y}$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	100	100	100	100	100
0	0	100	100	100	100	100
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0



3x3 image gradient filters

$$\frac{1}{3} \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix} \quad \frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$



Image gradients - Definition

Given function $f(x, y)$


Gradient vector $\nabla f(x, y) = \begin{bmatrix} \frac{\partial f(x, y)}{\partial x} \\ \frac{\partial f(x, y)}{\partial y} \end{bmatrix} = \begin{bmatrix} f_x \\ f_y \end{bmatrix}$

Gradient magnitude $|\nabla f(x, y)| = \sqrt{f_x^2 + f_y^2}$

Gradient direction $\theta = \tan^{-1} \frac{f_x}{f_y}$

Gradient of an image

- The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$



- The **gradient vector** points in the direction of most rapid increase in intensity
- The **gradient direction** is given by $\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$
 - how does this relate to the direction of the edge?
- The **edge strength** is given by the **gradient magnitude**

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Source: Steve Seitz

Image gradient


- Derivative of image wrt. x + Derivative wrt. y
→ Image gradient: magnitude and direction

- Gradient magnitude: gradient intensity for each pixel (*mostly used*)

$$|G| = \sqrt{(G_x^2 + G_y^2)} \approx |G_x| + |G_y|$$

- Gradient Direction: main direction of each pixel
 $\theta = \tan^{-1} G_y / G_x$

Example



- Which one is the gradient in the x-direction? How about y-direction?




CS231: Lecture 3 Juan Carlos Niebles and Ranjay Krishna
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

25

25

Edge detectors

Type of edges




Source : Gonzalez and Woods. Digital Image Processing 3ed. Prentice-Hall, 2008.

SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

26

26

Effects of noise



Source : Gonzalez and Woods. Digital Image Processing 3ed. Prentice-Hall, 2008.




SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

27

27

Effects of noise

- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal



Where is the edge?



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY


Source: S. Seitz

28

28

Effects of noise

- Solution: smoothing the image



Where is the edge?



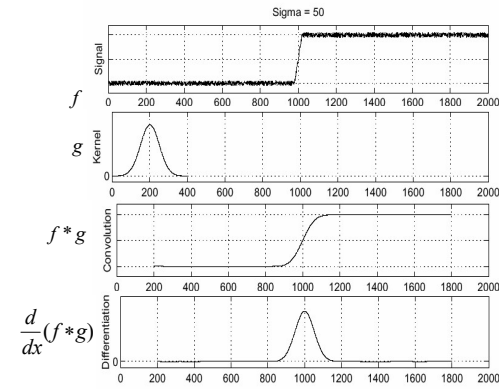
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Source: S. Seitz

29

29

Solution: smooth first



- To find edges, look for peaks in $\frac{d}{dx}(f * g)$

Source: S. Seitz



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

30


30

Derivative theorem of convolution

- This theorem gives us a very useful property:

$$\frac{d}{dx}(f * g) = \frac{df}{dx} * g = f * \frac{dg}{dx}$$

- This saves us one operation:



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Source: S. Seitz

31


31

Derivative theorem of convolution

- This theorem gives us a very useful property:

$$\frac{d}{dx}(f * g) = \frac{df}{dx} * g = f * \frac{dg}{dx}$$

- This saves us one operation:



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Source: S. Seitz

32

32

Discrete derivate in 2D image filters wrt. x and y

- Robert filter (the first approximation filter for image derivative - 1965)

$$\begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix} \quad \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$$

- Prewitt filter

$$\frac{1}{3} \times \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix} \quad \frac{1}{3} \times \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

- Sobel filter

$$\frac{1}{4} \times \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} \quad \frac{1}{4} \times \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$



Sobel Operator

- Gaussian Smoothing + differentiation

$$G_x = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$$

Gaussian smoothing
differentiation

$$G_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix} \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

→ Less sensible to noise



Prewitt Operator

- Mean smoothing + differentiation

$$G_x = \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$$


$$G_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 \end{bmatrix}$$

→ Less sensible to noise



Simple edge detector with 1st derivative

- Convolve the original image with 2 kernels to calculate approximations of the derivatives



- Compute the gradient magnitude



- Thesholding: choose edges to be the pixel above a threshold T



Simple edge detector with 1st derivative

Original image



Gradient magnitude using Sobel operator



Threshold T = 25



Threshold T = 60




SOURCE : Caroline Rougier, Traitement d'images (IFT2730), Univ. de Montréal.

37

37

Problems



- Poor localization (multiple adjacent pixels)
- Thresholding value favors certain directions over others
 - Can miss oblique edges more than horizontal or vertical edges → False negatives



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

38

Edge detector with 2nd derivative



Zero-crossing



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

39

39

Edge detector with 2nd derivative

- 2nd derivative with Laplacian filter:
 - convolution the image with one of 2 filters

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix} \quad \text{or} \quad \begin{bmatrix} 1 & 1 & 1 \\ 1 & -8 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

- Edge detection
 - Compute the 2nd derivative of the images
 - Find the zero-crossing pixels → edges



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

40


40

Edge detector with 2nd derivative

Image



Laplacian



- Single response
- Sensible to noise



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY


41

41

“Optimal” edge detector

• Criteria:

- **Good detection:** the optimal detector must minimize the probability of false positives (detecting spurious edges caused by noise), as well as that of false negatives (missing real edges)
- **Good localization:** the edges detected must be as close as possible to the true edges
- **Single response:** the detector must return one point only for each true edge point; that is, minimize the number of local maxima around the true edge



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

42

Canny detector

- This is probably the **most widely used** edge detector in computer vision
- Theoretical model: step-edges corrupted by additive Gaussian noise
- Canny has shown that the first derivative of the Gaussian closely approximates the operator that **optimizes the product of signal-to-noise ratio and localization**
- Optimal:
 - Detection: **weak edges detected**
 - Good location: **close to the real edges**
 - Unique response: **edge thickness = 1**

J. Canny, [A Computational Approach To Edge Detection](#), IEEE Trans. PAMI, 8:679-714, 1986.

SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

43

Canny detector: Steps

1) Apply a gaussian filter on the image


- Lowpass filter to **remove noise**

2) Compute the **gradient intensity** in the image

- Sobel filter in X and Y
- Compute the magnitude $|G| = |G_x| + |G_y|$

3) Compute image **gradient direction**

- Gradient direction $\theta = \arctan(G_y / G_x)$
- Round directions using multiples of 45°



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY


44

Canny detector: Steps

4) Non-maxima suppression

- If the gradient magnitude of a pixel (x,y) is inferior to the one of its 2 neighbors along the gradient direction

→ set this magnitude for (x,y) to zero




SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

45

Canny detector: Steps

5) Edge thresholding (hysteresis)


- Using two thresholds: a threshold high (S_h) and a threshold low (S_b)
- For each pixel in the gradient magnitude:
 - IF $\text{magnitude}(x,y) < S_b$, THEN set the pixel to zero (non-edge)
 - IF $\text{magnitude}(x,y) > S_h$, THEN the edge is an edge
 - IF $S_b \leq \text{magnitude}(x,y) \leq S_h$, THEN the pixel is an edge IF it is connected to another edge pixel



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

46

Hysteresis thresholding




SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Source: S. Seitz

47


Example: Canny detector



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

48

45 years of boundary detection



Source: Arbelaez, Maire, Fowlkes, and Malik. TPAMI 2011 (pdf)

Edge linking

Hough Transform
RANSAC
Edge closing


Hough transform

- The Hough transform (HT)
 - can be used to detect lines
 - was introduced in 1962 (Hough 1962) and first used to find lines in images a decade later (Duda 1972)
 - Goal: to find the location of lines in images.
- **Caveat:** Hough transform can *detect lines, circles* and other *structures* ONLY if their *parametric equation is known*
- It can give **robust detection** under *noise* and *partial occlusion*

Hough transform

- **Global** approach to detect continuous edges
 - *From the x-y plane to the parametric plane a-b*
- **x-y plane**
 - $y_i = a x_i + b$
 - an infinity of lines going through one (x_i, y_i) pair
 - one sole line for the (a,b) pair
- **a-b parametric plane**
 - $b = -x_i a + y_i$
 - one sole line for the (x_i, y_i) pair
 - an infinity of lines going through one (a,b) pair


x-y plane vs a-b plane



53

Line vs Points


All the points (x, y) on a line in the **x-y plane** are going through one sole point (a', b') in the **a-b parametric plane**



54

Main idea for the Hough transform

- Accumulation cells - Matrix (a, b)
- Build an **voting image**
 - each point is voting for a particular line
- The lines receiving more votes are kept



55

Computing the Hough transform


- We compute the **gradient** of the input image
 - *Sobel, Prewitt, Canny, ...*
- For each gradient point, we compute a line (a, b)
 - *Result is one line in the a-b plane for each pixel (x, y)*
- The maximum peaks in the a-b parametric plane show the lines with the maximum of points in the x-y plane
 - *The points indicating line crossing in the a-b plane show the real lines existing in the x-y plane*



56


Problem with the (a,b) space

- **Problem:** for a vertical line, $a=\infty$!
 - **Solution:** representing using polar coordinates (ρ, θ)
- $\rho = x \cos \theta + y \sin \theta$



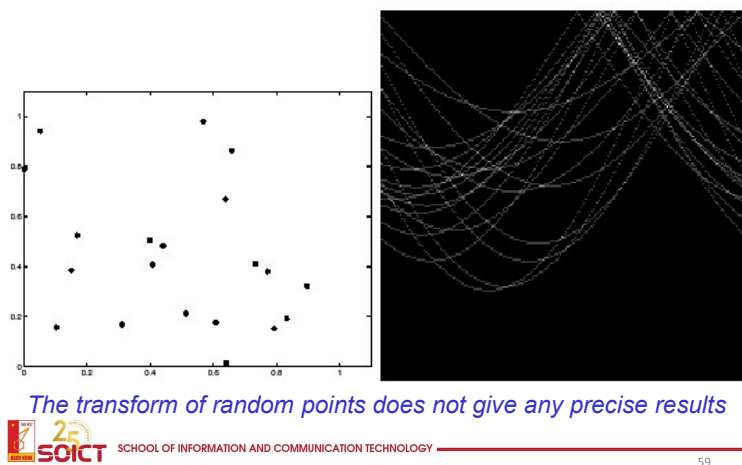
57

Example with 5 points



58


Hough transform (random points)



59


Hough transform (straight line)

The transform for aligned points result in a line detection



60

Example



61


RANSAC (RANdom Sample Consensus)

- A model fitting method:
 - A learning technique to estimate parameters of a model by random sampling of observed data
 - Used for :
 - Line detection
 - Correspondance problem (matching between 2 sets of features)
 - ...

62

Example: Line Fitting

- Why fit lines?
 - Many objects characterized by presence of straight lines



Slide credit: Kristen Grauman

63

Fitting as search in parametric space

- Choose a **parametric model** to represent a set of features
- Membership criterion is **not local**
 - Can't tell whether a point belongs to a given model just by looking at that point.
- Three main questions:
 - **What model** represents this set of features best?
 - **Which of several model instances** gets which feature?
 - How **many model instances** are there?
- Computational **complexity** is important
 - It is infeasible to examine every possible set of parameters and every possible combination of features

Source: L. Lazebnik

64

RANSAC [Fischler & Bolles 1981]

- **RAN**dom **SA**mple **C**onsensus
- Approach: we want to avoid the impact of outliers, so let's look for "inliers", and use only those
- Intuition: if an outlier is chosen to compute the current fit, then the resulting line won't have much support from rest of the points.



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide credit: Kristen Grauman

65

RANSAC [Fischler & Bolles 1981]

RANSAC loop:

1. Randomly **select a seed group** of points on which to base transformation estimate (e.g., a group of matches)
 2. **Compute transformation** from seed group
 3. Find **inliers** to this transformation
 4. If the number of inliers is **sufficiently large**, **re-compute least-squares estimate** of transformation on all of the inliers
- Keep the transformation with the largest number of inliers




SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide credit: Kristen Grauman

66

Difficulty of Line Fitting

- Extra edge points (clutter), multiple models:
 - Which points go with which line, if any?
- Only some parts of each line detected, and some parts are missing:
 - How to find a line that bridges missing evidence?
- Noise in measured edge points, orientations:
 - How to detect true underlying parameters?



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide credit: Kristen Grauman

67

Voting

- It's not feasible to check all combinations of features by fitting a model to each possible subset.
- Voting is a general technique where we let the features vote for all models that are compatible with it.
 - Cycle through features, cast votes for model parameters.
 - Look for model parameters that receive a lot of votes.
- Noise & clutter features will cast votes too, *but* typically their votes should be inconsistent with the majority of "good" features.
- Ok if some features not observed, as model can span multiple fragments.




SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide credit: Kristen Grauman

68

RANSAC Line Fitting Example

- Task: Estimate the best line
 - How many points do we need to estimate the line?




SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide credit: Kristen Grauman

69

RANSAC Line Fitting Example

- Task: Estimate the best line



Sample two points




SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide credit: Kristen Grauman

70

RANSAC Line Fitting Example

- Task: Estimate the best line



Fit a line to them




SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide credit: Kristen Grauman

71

RANSAC Line Fitting Example

- Task: Estimate the best line

Total number of points
within a threshold of
line.


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide credit: Kristen Grauman

72

RANSAC Line Fitting Example

- Task: Estimate the best line




SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide credit: Kristen Grauman

73

RANSAC Line Fitting Example

- Task: Estimate the best line




SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide credit: Kristen Grauman

74

RANSAC Line Fitting Example

- Task: Estimate the best line



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide credit: Kristen Grauman

75

Algorithm 15.4: RANSAC: fitting lines using random sample consensus

Determine:

- n — the smallest number of points required
- k — the number of iterations required
- t — the threshold used to identify a point that fits well
- d — the number of nearby points required to assert a model fits well

Until k iterations have occurred

Draw a sample of n points from the data uniformly and at random

Fit to that set of n points

For each data point outside the sample
Test the distance from the point to the line against t ; if the distance from the point to the line is less than t , the point is close

end

If there are d or more points close to the line then there is a good fit. Refit the line using all these points.

end

Use the best fit from this collection, using the fitting error as a criterion



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

76

RANSAC: How many samples?

- How many samples are needed?
 - Suppose w is fraction of inliers (points from line).
 - n points needed to define hypothesis (2 for lines)
 - k samples chosen.
 - Prob. that a single sample of n points is correct: w^n
 - Prob. that all k samples fail is: $(1 - w^n)^k$
- ⇒ Choose k high enough to keep this below desired failure rate.



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide credit: David Lowe

77

RANSAC: Computed k (p=0.99)

Sample size n	Proportion of outliers						
	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177




SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide credit: David Lowe

78

After RANSAC

- RANSAC divides data into inliers and outliers and yields estimate computed from minimal set of inliers.
- Improve this initial estimate with estimation over all inliers (e.g. with standard least-squares minimization).
- But this may change inliers, so alternate fitting with re-classification as inlier/outlier.



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide credit: David Lowe

79

RANSAC: Pros and Cons

- **Pros:**
 - General method suited for a wide range of model fitting problems
 - Easy to implement and easy to calculate its failure rate
- **Cons:**
 - Only handles a moderate percentage of outliers without cost blowing up
 - Many real problems have high rate of outliers (but sometimes selective choice of random subsets can help)
- A voting strategy, the Hough transform can handle high percentage of outliers



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

80